

Revanth Guthula

✉ revanthgwork@gmail.com | 📞 +91 90326 58749 | 🌐 github.com/rev2607 |
🌐 linkedin.com/in/rev2607 | 🌐 www.revanth.cloud

Summary

Final-year B.Tech student in AI & Data Science with strong skills in C++, Python, and system design. Passionate about scalable software development and applying AI/ML to solve real-world problems.

Skills

Languages: C++, Python, SQL

Systems & Tools: Linux, Docker, Git, Databases

Cloud Platforms: AWS, GCP, Azure, Vercel, Render

AI/ML: NumPy, Pandas, Scikit-learn, TensorFlow, PyTorch, FastAPI, LLMs (OpenAI, Ollama), LangChain

Work Experience

Ghar Ka Khana, Mumbai

Jul 2025 - Present

Software Developer Intern

- Designed and implemented a distributed, fault-tolerant web application for restaurant ordering using Next.js and Supabase - ensuring high availability and scalability for thousands of concurrent users.
- Optimized API performance and database queries, reducing response time by 30% under heavy load.
- Built authentication, secure API layers using Supabase & Twilio, ensuring data integrity and seamless user experience.
- Tech stack: HTML, CSS, TypeScript Next.js, React, Supabase, Twilio, AWS, Docker

Project Work

Mini RAG Q&A Bot with LangChain and OpenAI

[GitHub Link](#)

- Built an end-to-end RAG system for private datasets, delivering context-aware answers with sub-500ms latency.
- Productionized a lightweight API with observability, optimized embeddings/chunking, and CI-ready mock pipelines.
- Tech stack: Python, FastAPI, LangChain, OpenAI API, ChromaDB, Hugging Face tokenizers/embeddings, pytest

Redis-X (Redis-like In-Memory Database)

[GitHub Link](#)

- Designed a scalable, high-performance, event-driven server with a custom binary protocol; optimized concurrent I/O and pipelining, reducing p95 latency and increasing throughput.
- Built progressive-rehash hash map, AVL ZSETs, and min-heap TTL; implemented fault-tolerant expiration and monitoring, enabling $O(1)/O(\log n)$ ops and cost-efficient memory usage.
- Tech stack: C++17, Linux, TCP sockets, pthreads, Hash map, AVL tree, Heap, Python

DataPulse — Scalable Data Health Monitoring System

[GitHub Link](#)

- Built to ensure data quality in enterprise-scale pipeline analyzing 10M+ records for missing values, outliers, and drift.
- Used asynchronous task execution and containerization to reduce data validation time by 40%.
- Tech stack: FastAPI, Python, Scikit-learn, React, Docker

Education

KL University Hyderabad

Aug 2022 - May 2026 (Expected)

Bachelor's of Technology in Artificial Intelligence and Data Science

CGPA: 7.76/10

Relevant Coursework: Data Structures & Algorithms, Operating Systems, Computer Networks, Database Management Systems, Machine Learning, Deep Learning, Big Data Analytics, Probability & Statistics, Data Science & Visualization

Achievements & Certifications

- Published "An Enhanced SVM Model Optimized with Minimum Bayes Error Rate for Mental Disorder Detection" — Springer ICT4SD 2025
- Google Cloud Digital Leader — Certified
- Microsoft Azure AI Engineer Associate — Certified
- NXP AIM India 2024 (Regionals)** – Built an obstacle-avoiding robot with ROS2 lane detection, applying real-time decision-making and distributed robotics concepts.